

Introduction to Resampling in R

DQuinn

Friday, November 06, 2015

Ecological studies often suffer from limited resource availability, including time, money, and people power. Oversampling not only taxes these resources needlessly, but may also negatively influence the system being studied; for example, through lethal collections of organisms or unnecessary handling of non-target species. Projects relying on data collection over multiple field seasons may benefit from resampling algorithms that can identify efficient collection protocols if oversampling is known or thought to have occurred.

We'll be using a mock data set of tagged and non-tagged eels captured in a small lake to estimate population size, then build a resampling algorithm to explore the sensitivity of using varying numbers of traps and sampling events.

- 50 American eels were captured from a small lake, marked with an external tag, and released back into the lake
- 15 traps were checked 20 times over a season
- number of tagged and non-tagged eels in each trap was recorded

Load Required Packages

```
library(dplyr)
library(FSA)
library(ggplot2)
library(digest)
library(lubridate)
```

You can find the dataset here:

https://github.com/DanielleQuinn/RLessons/blob/master/SensitivityAnalysis/recapture_data.csv

```
df.data <- read.csv("recapture_data.csv")
```

First, take a look at the data frame `df.data`

```
head(df.data)
```

	X	trap	check	present	tag
1	1	1	1	1	1
2	2	1	1	1	0
3	3	1	1	1	0
4	4	1	1	1	0
5	5	1	1	1	0
6	6	1	1	1	0

```
str(df.data)
```

```
'data.frame': 150036 obs. of 5 variables:
 $ X      : int  1 2 3 4 5 6 7 8 9 10 ...
 $ trap   : int  1 1 1 1 1 1 1 1 1 1 ...
 $ check  : int  1 1 1 1 1 1 1 1 1 2 ...
 $ present: int  1 1 1 1 1 1 1 1 1 1 ...
 $ tag    : int  1 0 0 0 0 0 0 0 0 1 ...
```

Baseline Population Size Estimate

We'll be using the `mrClosed()` function from the `FSA` package to estimate the population size using the Schnabel method. Without getting into too much detail, the Schnabel method is used when we have multiple sampling occasions and assumes a closed population. The function requires specific input, including:

- `n`: the number of captured animals
- `m`: the number of recaptured marked animals
- `M`: the number of extant marked animals prior to the sample

Step 1: Build cmr data frame

```
df.cmr<-data.frame(df.data%>%
  group_by(check)%>%
  summarise(date=unique(check),n=sum(present), m=sum(tag),
M=50)%>%
  select(n,m,M))
```

Step 2: Estimate population size

```
pop.est<-summary(with(df.cmr,mrClosed(n=n,m=m,M=M,method="Schnabel"))) #From
FSA package
pop.low<-
stats::confint(with(df.cmr,mrClosed(n=n,m=m,M=M,method="Schnabel")))[1] #
Confidence intervals #
pop.high<-
stats::confint(with(df.cmr,mrClosed(n=n,m=m,M=M,method="Schnabel")))[2] #
Confidence intervals #
baseline<-data.frame(pop.est, pop.low, pop.high)
baseline
```

	N	pop.low	pop.high
1	152	151	154

We see that using all of the data available (50 traps each checked 20 times), our population estimate is 152, with confidence intervals of 151 and 154.

Because the lake is very small, we're fairly certain that oversampling may have occurred, and our goal is to design a sampling protocol for the next year that maximizes the efficiency of our sampling effort. We want to know the minimum number of traps and the minimum number of checks that we can do that will result in a realistic estimate of population size.

Sample 10 Random Traps

```
all_traps<-unique(df.data$trap) # A vector of all available traps
use_traps<-sample(all_traps, 10) # Choose 10 random traps
```

Subset the data to include only those traps, then create `df.cmr` and use `mrClosed()` to estimate population size.

```
my_subset<-data.frame(df.data%>%
  filter(trap %in% use_traps))

df.cmr<-data.frame(my_subset%>%
  group_by(check)%>%
  summarise(date=unique(check), n=sum(present), m=sum(tag),
M=50)%>%
  select(n,m,M))

pop.est<-summary(with(df.cmr, mrClosed(n=n,m=m,M=M,method="Schnabel"))) #From
FSA package
pop.low<-
stats::confint(with(df.cmr, mrClosed(n=n,m=m,M=M,method="Schnabel")))[1] #
Confidence intervals #
pop.high<-
stats::confint(with(df.cmr, mrClosed(n=n,m=m,M=M,method="Schnabel")))[2] #
Confidence intervals #
results<-data.frame(pop.est, pop.low, pop.high)
results

      N pop.low pop.high
1 163      159      167
```

Sample X Random Traps

Repeat this analysis using varying numbers of traps and see how our results vary from the baseline population estimates. We are going to resample from 10 to 50 random traps and apply to analysis for each subset.

Step 1: Build a function that creates the `cmr` data frame and produces the population estimate and confidence intervals

```
my_cmr<-function(my_subset)
{
  # Build cmr data frame
  df.cmr<-data.frame(my_subset%>%
    group_by(check)%>%
    summarise(date=unique(check), n=sum(present),
m=sum(tag), M=50)%>%
    select(n,m,M))

  # Estimate Population Size
  pop.est<-summary(with(df.cmr, mrClosed(n=n,m=m,M=M,method="Schnabel")))
#From FSA package
pop.low<-
```

```
stats::confint(with(df.cmr,mrClosed(n=n,m=m,M=M,method="Schnabel")))[1] #
Confidence intervals #
  pop.high<-
stats::confint(with(df.cmr,mrClosed(n=n,m=m,M=M,method="Schnabel")))[2] #
Confidence intervals #
}
```

Step 2: Set up a data frame to handle the output of the analyses We want a data frame that has information about how many traps were used, and the results of our analysis. In this case, we're going to have 41 sets of output.

```
results<-data.frame(traps=c(10:50), pop.est=NA, pop.low=NA, pop.high=NA)
results
```

	traps	pop.est	pop.low	pop.high
1	10	NA	NA	NA
2	11	NA	NA	NA
3	12	NA	NA	NA
4	13	NA	NA	NA
5	14	NA	NA	NA
6	15	NA	NA	NA
7	16	NA	NA	NA
8	17	NA	NA	NA
9	18	NA	NA	NA
10	19	NA	NA	NA
11	20	NA	NA	NA
12	21	NA	NA	NA
13	22	NA	NA	NA
14	23	NA	NA	NA
15	24	NA	NA	NA
16	25	NA	NA	NA
17	26	NA	NA	NA
18	27	NA	NA	NA
19	28	NA	NA	NA
20	29	NA	NA	NA
21	30	NA	NA	NA
22	31	NA	NA	NA
23	32	NA	NA	NA
24	33	NA	NA	NA
25	34	NA	NA	NA
26	35	NA	NA	NA
27	36	NA	NA	NA
28	37	NA	NA	NA
29	38	NA	NA	NA
30	39	NA	NA	NA
31	40	NA	NA	NA
32	41	NA	NA	NA
33	42	NA	NA	NA
34	43	NA	NA	NA
35	44	NA	NA	NA

36	45	NA	NA	NA
37	46	NA	NA	NA
38	47	NA	NA	NA
39	48	NA	NA	NA
40	49	NA	NA	NA
41	50	NA	NA	NA

Step 3: Resample the data and fill in the results data frame


```
for(i in 1:nrow(results)) # For each row in our results data frame
{
  use_traps<-sample(all_traps, results$traps[i]) # Sample x number of random
traps
  my_subset<-data.frame(df.data%>%filter(trap %in% use_traps)) # Subset the
data to only include those traps
  my_cmr(my_subset) # Apply our my_cmr function to the subset
  results$pop.est[i]<-pop.est # Fill in our results
  results$pop.low[i]<-pop.low
  results$pop.high[i]<-pop.high
}
results
```

	traps	pop.est	pop.low	pop.high
1	10	145	142	148
2	11	152	149	155
3	12	143	141	145
4	13	150	147	153
5	14	160	157	163
6	15	151	148	154
7	16	148	146	151
8	17	157	154	160
9	18	152	149	154
10	19	155	152	157
11	20	152	150	154
12	21	151	149	153
13	22	152	149	154
14	23	152	150	154
15	24	147	145	149
16	25	159	157	162
17	26	155	153	158
18	27	156	154	158
19	28	150	148	152
20	29	151	149	153
21	30	153	151	155
22	31	156	154	158
23	32	152	150	154
24	33	154	152	156
25	34	153	151	155
26	35	151	150	153
27	36	153	151	154
28	37	151	150	153

29	38	151	149	152
30	39	154	153	156
31	40	151	149	152
32	41	151	149	152
33	42	153	152	155
34	43	155	153	156
35	44	150	149	152
36	45	154	152	155
37	46	151	150	152
38	47	154	153	156
39	48	154	152	155
40	49	152	151	153
41	50	152	151	154

Step 4: Visualize results

```
ggplot(results)+
  geom_line(aes(x=traps, y=pop.est))+
  geom_line(aes(x=traps, y=pop.low), linetype='dashed')+
  geom_line(aes(x=traps, y=pop.high), linetype='dashed')+
  geom_hline(yint=baseline$N, col='green4')+
  geom_hline(yint=baseline$pop.low, col='green4', linetype='dashed')+
  geom_hline(yint=baseline$pop.high, col='green4', linetype='dashed')+
  theme_bw(16)+xlab("Number of Traps")+ylab("Population Estimate")
```



Because random traps are selected each time, your results will look slightly different. However, you should see the same amount of variation in population estimates. Why is

there so much variation? It's because we're only randomly selecting each number of traps a single time. To really get an idea of the overall sensitivity of the population estimate to the number of traps used, we'll repeat this resampling algorithm many times.

Bootstrapping

Repeat this algorithm 25 times, but instead of varying the number of traps to be 1 to 50, we'll limit those options to be 10, 15, 20, ...50.

Step 1: Build a function - already done!

Step 2: Set up a data frame to handle the output of the analyses In this case, we're going to have 9 (different numbers of traps to select) x 25 (repeats) outputs:

```
results<-data.frame(traps=rep(seq(from=10, to=50, by=5),25), pop.est=NA,
pop.low=NA, pop.high=NA)
head(results)
```

	traps	pop.est	pop.low	pop.high
1	10	NA	NA	NA
2	15	NA	NA	NA
3	20	NA	NA	NA
4	25	NA	NA	NA
5	30	NA	NA	NA
6	35	NA	NA	NA

Step 3: Resample and fill in the results data frame In addition, we'll set up a means of keeping track of how long the resampling takes, and how long each iteration takes to complete.

```
starttime=Sys.time() # What time does the resampling begin?
for(i in 1:nrow(results))
{
  use_traps<-sample(all_traps, results$traps[i])
  my_subset<-data.frame(df.data%>%filter(trap %in% use_traps))
  my_cmr(my_subset)
  results$pop.est[i]<-pop.est
  results$pop.low[i]<-pop.low
  results$pop.high[i]<-pop.high
}
totaltime<-difftime(Sys.time(),starttime, unit="secs") #How Long did it take?
totaltime


Time difference of 5.752329 secs

paste(round(as.numeric(totaltime)/nrow(results),2),"seconds per iteration") #
How Long did each iteration take?

[1] "0.03 seconds per iteration"
```

Step 4: Visualize results

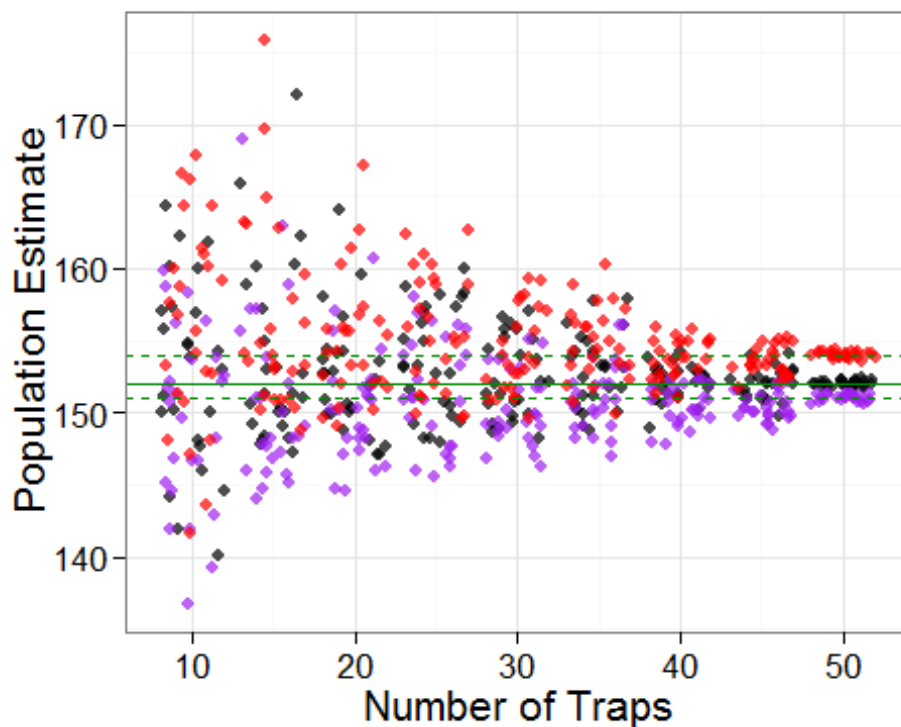
```
ggplot(results)+
  geom_line(aes(x=traps, y=pop.est))+
  geom_line(aes(x=traps, y=pop.low), linetype='dashed')+
  geom_line(aes(x=traps, y=pop.high), linetype='dashed')+
  geom_hline(yint=baseline$N, col='green4')+
  geom_hline(yint=baseline$pop.low, col='green4',linetype='dashed')+
  geom_hline(yint=baseline$pop.high, col='green4',linetype='dashed')+
  theme_bw(16)+xlab("Number of Traps")+ylab("Population Estimate")
```



Why do we have jagged lines? It's because each value of traps has 25 points associated with it. There are different ways to deal with this.

Option 1: Don't use a line to represent the resampled data.

```
ggplot(results)+
  geom_jitter(aes(x=traps, y=pop.est), alpha=0.7)+
  geom_jitter(aes(x=traps, y=pop.low), col='purple', alpha=0.7)+
  geom_jitter(aes(x=traps, y=pop.high), col='red', alpha=0.7)+
  geom_hline(yint=baseline$N, col='green4')+
  geom_hline(yint=baseline$pop.low, col='green4',linetype='dashed')+
  geom_hline(yint=baseline$pop.high, col='green4',linetype='dashed')+
  theme_bw(16)+xlab("Number of Traps")+ylab("Population Estimate")
```

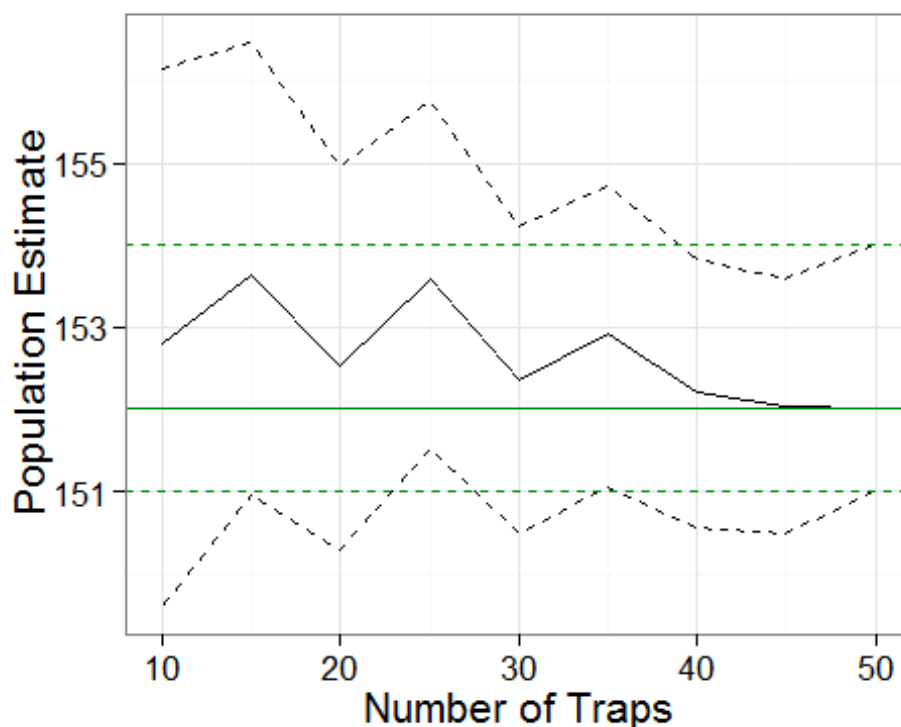



You can start to see that when few traps are used, the population estimates are much more variable than if many traps are used.

Option 2: Summarize the data.

```
results.sum<-data.frame(results%>%
  group_by(traps)%>%
  summarise(mean.pop=mean(pop.est), mean.low=mean(pop.low),
    mean.high=mean(pop.high)))

ggplot(results.sum)+
  geom_line(aes(x=traps, y=mean.pop))+
  geom_line(aes(x=traps, y=mean.low), linetype='dashed')+
  geom_line(aes(x=traps, y=mean.high), linetype='dashed')+
  geom_hline(yint=baseline$N, col='green4')+
  geom_hline(yint=baseline$pop.low, col='green4', linetype='dashed')+
  geom_hline(yint=baseline$pop.high, col='green4', linetype='dashed')+
  theme_bw(16)+xlab("Number of Traps")+ylab("Population Estimate")
```



More Bootstrapping

Recall that each trap was checked 20 times. Imagine all of the time and effort that would be saved if we could sample traps fewer times and still be confident in our population estimates! In addition to exploring the sensitivity of the number of traps, let's add another level to this analysis and include a varying number of checks. We'll say that we want to see what would happen if 10 to 20 checks were done.

```
all_checks<-unique(df.data$check)
```

Because we've been keeping track of how long each iteration takes (approx. 0.03 sec, depending on the processing power of your computer), we know that we can easily do more iterations within a reasonable time frame. We'll do 40 repeats this time.

Step 1: Build a function - already done!

Step 2: Set up a data frame to handle the output of the analyses In this case, we're going to have 9 (different numbers of traps to select) x 11 (different numbers of checks to select) x 40 (repeats) outputs. Using `expand.grid()` will be useful here; it creates a dataframe of every combination of each argument you give it.

```
results<-expand.grid(traps=seq(from=10, to=50, by=5), checks=c(10:20),
repeats=c(1:40))
nrow(results) # How many iterations?
[1] 3960
nrow(results)*0.03/60 # It should take roughly 2 minutes
```

```
[1] 1.98
```

```
# We can get rid of the repeats column
```

```
results<-results[, -3]
```

```
# And need to add the pop.est, pop.low, and pop.high columns
```

```
results$pop.est<-NA
```

```
results$pop.low<-NA
```

```
results$pop.high<-NA
```

Step 3: Resample and fill in the results data frame

```
starttime=Sys.time()
```

```
for(i in 1:nrow(results))
```

```
{
```

```
# print(paste(i/nrow(results)*100, "% Complete")) # This is useful for  
keeping track of how many iterations have been completed
```

```
use_traps<-sample(all_traps, results$traps[i])
```

```
use_checks<-sample(all_checks, results$checks[i])
```

```
my_subset<-data.frame(df.data%>%filter(trap %in% use_traps)%>%filter(check  
%in% use_checks))
```

```
my_cmr(my_subset)
```

```
results$pop.est[i]<-pop.est
```

```
results$pop.low[i]<-pop.low
```

```
results$pop.high[i]<-pop.high
```

```
}
```

```
totaltime<-difftime(Sys.time(), starttime, unit="secs")
```

```
totaltime
```

```
Time difference of 109.2773 secs
```

```
paste(round(as.numeric(totaltime)/nrow(results), 2), "seconds per iteration")
```

```
[1] "0.03 seconds per iteration"
```

Step 4: Visualize results

```
results.sum<-data.frame(results%>%
```

```
group_by(traps, checks)%>%
```

```
summarise(mean.pop=mean(pop.est),
```

```
mean.low=mean(pop.low), mean.high=mean(pop.high))
```

```
ggplot(results.sum)+
```

```
geom_line(aes(x=traps, y=mean.pop))+
```

```
geom_line(aes(x=traps, y=mean.low), linetype='dashed')+
```

```
geom_line(aes(x=traps, y=mean.high), linetype='dashed')+
```


```
geom_hline(yint=baseline$N, col='green4')+
```

```
geom_hline(yint=baseline$pop.low, col='green4', linetype='dashed')+
```

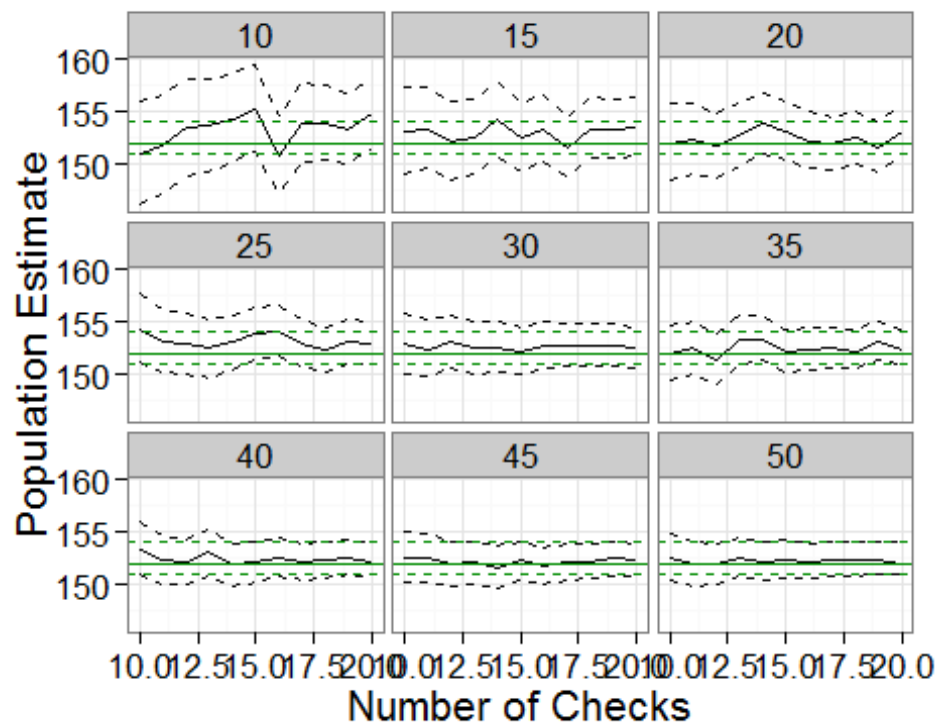
```
geom_hline(yint=baseline$pop.high, col='green4', linetype='dashed')+
```

```
theme_bw(16)+xlab("Number of Traps")+ylab("Population Estimate")+
```

```
facet_wrap(~checks)
```



```
ggplot(results.sum)+
  geom_line(aes(x=checks, y=mean.pop))+
  geom_line(aes(x=checks, y=mean.low), linetype='dashed')+
  geom_line(aes(x=checks, y=mean.high), linetype='dashed')+
  geom_hline(yint=baseline$N, col='green4')+
  geom_hline(yint=baseline$pop.low, col='green4', linetype='dashed')+
  geom_hline(yint=baseline$pop.high, col='green4', linetype='dashed')+
  theme_bw(16)+xlab("Number of Checks")+ylab("Population Estimate")+
  facet_wrap(~traps)
```



You can use these plots to start identifying the number of traps and the number of checks that result in population estimates that you consider to be acceptable, and from there you can design more efficient sampling protocols for next year!